

# Deep Learning

## Getting Started



Jeremy Pinto



jeremy@jerpint.io

<http://>

# What this talk is about

- > Overview of deep learning
- > Examples of how it can be used
  - Image processing
  - Text analysis
- > Tips on getting started

# Machine Learning

- > Teach machines to learn abstract concepts.
- > Example: what is a chair?



# Machine Learning

> What about these?




# Semantics

> Typical machine learning problem:

- **X**: set of examples
- **y**: target

Find  $f(\mathbf{X}) \approx \mathbf{y}$



# Deep Learning

# Example

**x**: Input (image)

**y**: Target (hot dog | not hot dog)

$f(\mathbf{x}) = \mathbf{y}'$  = Prediction

$L(\mathbf{y}', \mathbf{y})$ : Loss (error function)



# Example

$$f(\text{img}) = \begin{array}{ll} \text{Not hotdog!} & p = 0.59 \\ \text{Hotdog!} & p = 0.41 \end{array}$$

$$f(\text{img}) = \begin{array}{ll} \text{Not hotdog!} & p = 0.99 \\ \text{Hotdog!} & p = 0.01 \end{array}$$

Train on  
examples

Hotdog!




Not hotdog!







# Deep learning system



# Deep learning system

Input =  $\mathbf{X}$  



Black Box =  $f(\mathbf{X})$




Prediction =  $\mathbf{y}'$



# Deep learning system

- > Linear Algebra
- > Calculus
- > Non-Linearities



Black Box =  $f(\mathbf{X})$



“Linear algebra is one of the fundamental mathematical disciplines necessary to understanding deep learning.”  
– The Deep Learning Book

<http://www.deeplearningbook.org>

# Deep Learning – Optimization

Goal : Minimize the error of our black box

Given:  $f(\mathbf{X}) = \mathbf{y}'$  = prediction

Error:  $L(\mathbf{y}', \mathbf{y})$  = how good is our prediction?


Given our error, tune the parameters in the model

# Recall: Minimization

Calculus: given a function  $L(x)$ , find its minimum

$$\frac{\partial L}{\partial x} = 0$$


$$\frac{\partial^2 L}{\partial x^2} > 0$$





x

# Recall: Minimization

Deep Learning – The loss function and its derivative is only approximated locally.



Control  $x$




# Recall: Minimization

Deep Learning – The loss function and its derivative is only approximated locally.



Control  $x$




Gradient Descent

$$x' = x - \epsilon \nabla_x f(x)$$

<http://www.deeplearningbook.org>

# Deep learning system

- > Linear Algebra
- > Calculus
- > Non-Linearities



Black Box =  $f()$





# Image Processing

# MNIST - Hello (Deep Learning) World

2 2 2 2 2 2  
9 9 9 9 9 9  
1 1 1 1 1 1  
1 1 1 1 1 1  
3 3 3 3 3 3  
5 5 3 5 3 3  
2 2 2 2 3 2  
7 7 7 7 7 7  
0 0 0 0 0 0  
0 0 0 0 0 0  
4 4 4 4 4 9  
6 6 6 6 6 6  
5 5 5 5 5 5  
9 9 9 9 4 9  
3 3 3 3 3 3  
7 7 7 7 7 7  
0 0 0 0 0 0  
6 6 6 6 6 6  
9 9 7 4 9 9  
8 8 8 8 8 8


## Dataset

A 10x10 grid of black and white squares. The number '88' is formed by black squares on a white background. The first '8' is located in the first two columns, and the second '8' is located in the last two columns. Each '8' is composed of three stacked loops.

28 x 28  
784 pixels



[illegible]

# Multi Layer Perceptron (Fully Connected)



Digit  
probability

# Multi Layer Perceptron (Fully Connected)



# Convolution


$3_0$	$3_1$	$2_2$	1	0
$0_2$	$0_2$	$1_0$	3	1
$3_0$	$1_1$	$2_2$	2	3
2	0	0	2	2
2	0	0	0	1

 $\otimes$ 
$$\begin{pmatrix} 0 & 1 & 2 \\ 2 & 2 & 0 \\ 0 & 1 & 2 \end{pmatrix}$$
 $=$ 


12	12	17
10	17	19
9	6	14

[http://deeplearning.net/software/theano/tutorial/conv\\_arithmetic.html](http://deeplearning.net/software/theano/tutorial/conv_arithmetic.html)

# LeNet (convolutional neural network)




## Convolutional Neural Nets (CNNs): 1989



LeNet: a layered model composed of convolution and subsampling operations followed by a holistic representation and ultimately a classifier for handwritten digits. [ LeNet ]

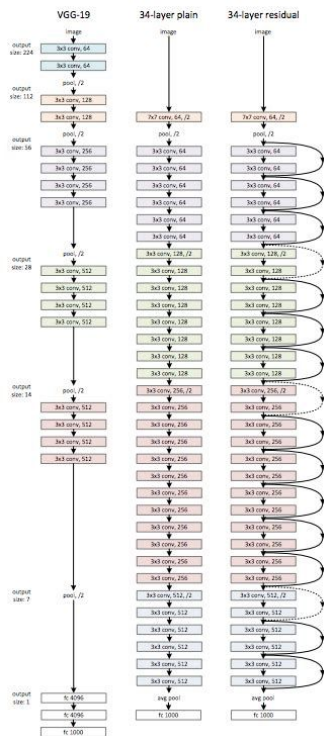
# AlexNet – ImageNet Winner 2012

## ImageNet Dataset




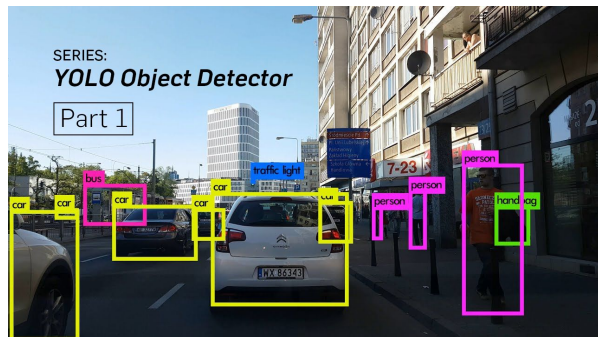
Russakovsky, O., Deng, J., Su, H., Krause, J., Satheesh, S., Ma, S., ... & Fei-Fei, L. (2015). [Imagenet large scale visual recognition challenge](#). *arXiv preprint arXiv:1409.0575*. [\[web\]](#)

# Deep Learning Revolution







# Deep Learning Applications (images)



Mask R-CNN




# Deep Learning Applications (images)



# Natural Language Processing

# Word Embeddings (word2vec, Glove etc.)

- > Learn a high-dimensional vector representation for words



<https://papers.nips.cc/paper/5021-distributed-representations-of-words-and-phrases-and-their-compositionality.pdf>

<http://mccormickml.com/2016/04/19/word2vec-tutorial-the-skip-gram-model/>

# NLP – Char–RNN

VIOLA:

Why, Salisbury must find his flesh and thought  
That which I am not apt, not a man and in fire,  
To show the reining of the raven and the wars  
To grace my hand reproach within, and not a fair are hand,  
That Caesar and my goodly father's world;  
When I was heaven of presence and our fleets,  
We spare with hours, but cut thy council I am great,  
Murdered and by thy master's ready there  
My power to give thee but so much as hell:  
Some service in the noble bondman here,  
Would show him to her wine.

KING LEAR:

O, if you were a feeble sight, the courtesy of your law,  
Your sight and several breath, will wear the gods  
With his heads, and my hands are wonder'd at the deeds,  
So drop upon your lordship's head, and your opinion  
Shall be against your honour.

CINNA


No, by no means.

METELLUS CIMBER

O, let us have him, for his silver hairs  
Will purchase us a good opinion  
And buy men's voices to commend our deeds:  
It shall be said, his judgment ruled our  
hands;  
Our youths and wildness shall no whit  
appear,  
But all be buried in his gravity.

<https://karpathy.github.io/2015/05/21/rnn-effectiveness/>

# Recurrent Neural Networks



AI TIP: TO DEVELOP A COMPUTER WITH THE INTELLIGENCE OF A SIX-YEAR-OLD CHILD, START WITH ONE AS SMART AS AN ADULT AND LET ME TEACH IT STUFF.

# Getting Started


# Python



*matplotlib*

pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$






# Deep Learning Libraries

> Start with tutorials

theano



> Focus on inputs and outputs to common architectures

PYTORCH




# Explicitly define your problem

**X:** What is your input?

**y:** What is your target?

Define proper evaluation  
metrics

What architecture should you  
use?



# Split your data

You should have 3 sets of data


- > Train (~80%)
- > Validation (~10%)
- > Testing (~10%)



This will avoid overfitting your data

# Overfitting


- > Can happen when your data is noisy or if little training data is available
- > Signs of overfit
  - 100% accuracy?
  - test set and train set too similar?
  - Data leak? Bias in the data?



# Fine-Tuning

- > Use an existing model as a starting point for your experiment
  - Great for smaller datasets
  - Works well in practice
  - Allows you to save a lot of implementation time!
- > Example - Use a model trained on ImageNet to classify new classes
  - Filters learned are already generalized enough
  - Takes a lot less time to converge

Fine-tuning Pretrained Network



# Do's

- > Look for pre-trained models
- > Look for similar datasets
- > Adapt code to suit your needs
- > Ask questions
- > Read A LOT!
  - Lots of great blog posts!
  - Understand concepts first, understand details after

# Don't's

- > Reimplement papers from scratch  
(Right away)
  - Great for learning
  - Time consuming
  - Lots of debugging
  - implementation details aren't always clear
- > Train from scratch
  - Learning from scratch is hard and resource intensive
  - Fine-tune instead!

# GPU vs CPU





# Computing Resources

- > Most libraries can run on CPU
  - Inference (making a prediction) is fine on CPU
- > GPU is ideal for training larger models
- > GPUs are expensive and need the right drivers (CUDA, CUDNN)
- > Cloud alternatives : Amazon, Azure, Google Cloud etc.



**nVIDIA**



Microsoft Azure



Google Cloud

# Learn the jargon

Loss, epoch, splits, learning rate,  
accuracy, dropout, early-stopping,  
layers, convolution, RNN, LSTM, MLP,  
Fully-connected etc.

# Practical Considerations

- > How large are your models?
- > Can they work offline? (edge computing)
- > How do they scale?
- > Is it for real-time applications?
- > How will you deploy?
- > How does it compare to simpler models?

# Use the right tools



# Questions?



Jeremy Pinto



jeremy@jerpint.io

<http://>